

A Trillion Vectors in the Smallest Job Class [Scalable Data Science]

Building and measuring a sharded vector index on a shared research cluster with one CPU and two gigabytes per job

Andrew H. Bond
San José State University
San José, California, USA
andrew.bond@sjsu.edu

ABSTRACT

A sharded vector index of one trillion 4-bit quantized rows, 24 TB on 500 block volumes, was built and then measured exactly on the National Research Platform, a Kubernetes cluster shared by many research groups, with no reserved capacity. Every job in the build and the measurement requested one CPU and two gigabytes of memory, the smallest class the cluster offers. Three design choices made that possible. The corpus is defined by a seed and regenerated where it is needed, so the build moved no corpus bytes and a lost volume was rebuilt from its seed in under four hours. Every job is idempotent on its volume, so a failed or duplicated job costs a retry and never a wrong result. A pool of per-server jobs, submitted through a controller that yields when the cluster is busy, recovers from failures by a small number of explicit rules rather than by operator attention. The build of 500 servers took seven weeks of background capacity, and the exact scan and five routed passes over the whole index took four days, 1349 CPU-hours, 1502 submissions and 1350 completions, with six pods replaced by hand. The recall this measured is reported in a companion paper. Here we describe the index layout, the corpus generator, the job shape, the pool and its rules, what failed and how often, and what the whole run cost.

PVLDB Reference Format:

Andrew H. Bond. A Trillion Vectors in the Smallest Job Class [Scalable Data Science]. PVLDB, 20(1): XXX-XXX, 2027.
doi:XX.XX/XXX.XX

PVLDB Artifact Availability:

The source code, data, and/or other artifacts have been made available at <https://github.com/ahb-sjsu/turboquant-pro/tree/master/benchmarks/fleet>.

1 INTRODUCTION

Measuring a vector index at a trillion rows needs a trillion rows in one place and an exact scan of all of them. Few groups own that much capacity. Research clusters that federate many sites have it in aggregate, but they share it, reserve nothing for any one user, and expect jobs that come and go without supervision. This paper describes how a trillion-row index was built and scanned exactly




Figure 1: The fleet. One bootstrap shard fixes the projection basis and the 2048-cell coarse quantizer for every shard on every server, so scores are comparable across servers and the global top-10 is the exact merge of per-server top-10s. Each phase is one job per server, and each job reads only its own volume and the shared volume.

on such a cluster, the National Research Platform’s Nautilus [12], with every job in the smallest resource class the platform offers.

The run exists to answer a measurement question, whether routed search over a sharded index loses anything the exact scan finds as the index grows to 10^{12} rows. The companion paper [3] reports that answer. This paper is about the machinery. It makes three points. A corpus defined by a seed turns a 24 TB build into computation with no data movement and turns the loss of a volume into a rebuild. A scan over memory-mapped shards fits one CPU and two gigabytes once three choices are made about queries, open shards and block size. A pool of idempotent jobs with explicit rules for failure, waiting and deferral carries a four-day run of 1502 submissions through the ordinary failures of a shared cluster. Section 10 sets each design choice against the choice a first implementation would make and gives what it measured in the run (Table 1), and states what we would change.

2 THE INDEX IS 500 SERVERS OF TWO BILLION ROWS

Index. Each row is a 32-dimensional vector projected to 24 dimensions by one basis shared by the whole fleet, quantized to 4 bits per dimension, and assigned to one of 2048 inverted-file cells by one shared coarse quantizer. Both the basis and the quantizer

This work is licensed under the Creative Commons BY-NC-ND 4.0 International License. Visit <https://creativecommons.org/licenses/by-nc-nd/4.0/> to view a copy of this license. For any use beyond those covered by this license, obtain permission by emailing info@vldb.org. Copyright is held by the owner/author(s). Publication rights licensed to the VLDB Endowment.
Proceedings of the VLDB Endowment, Vol. 20, No. 1 ISSN 2150-8097.
doi:XX.XX/XXX.XX

Table 1: The design choices of this paper, each beside the choice a first implementation would make, what that choice costs at 10^{12} rows, and what the design measured in the run.

obvious choice	its cost at 10^{12} rows	design	measured consequence
store the float corpus and read it	place and read 128 TB	corpus defined by a seed, regenerated by each job	no corpus rows moved, only code, the bootstrap and small results (Section 3)
build a shard in memory	a five-million-row block is 640 MB before codes and temporaries	stream row bands, flush and drop each band	build in one CPU and two gigabytes, 292 MiB anonymous (Table 4)
large jobs per server	scarce capacity, hard to place beside other groups	one CPU and two gigabytes per job	exact scan 0.88 cores and 1242 MiB mean, twenty in flight (Figure 3)
default 128 open shards	5.8 GB of row-id and tombstone arrays	at most two open, eight for the routed pass	exact-scan peak 1469 MiB within two gigabytes
write partials in place	a retry can meet a half-written result	temporary name and rename, exit if the output exists	174 re-issued jobs over both runs, each safe to repeat (Table 5)
resubmit a job that is not yet visible	a duplicate behind every deferred job	wait ninety minutes for the controller	the one deferral of the first run absorbed without a duplicate
back up the index	a second copy of 24 TB	rebuild a lost volume from its seed	six volumes rebuilt, 2 h 32 min to 3 h 49 min each, no backup stored

are fitted once, on a bootstrap shard, and read by every build job. A shard holds five million rows. A server holds 400 shards, two billion rows, on one block volume. The on-disk row is 12 bytes of codes, 8 bytes of per-row norms used by the asymmetric distance computation, and 4 bytes of cell membership. Every server reports 48 009 656 624 index bytes for its two billion rows, 24.0 bytes per row, and the fleet holds 24.0 TB of index on 500 volumes. Scoring is asymmetric distance computation over the codes, in row blocks, from memory-mapped shards. Because the basis and the coarse quantizer are shared, scores from different shards are comparable and a global top- k is the exact merge of per-shard top- k s. That property is what makes an exact measurement of the whole index possible, and Figure 1 shows the arrangement.

Scale. The same layout has been built and measured at 10^8 rows on 4 servers, 10^9 on 4, 10^{10} on 8, 10^{11} on 50 and 10^{12} on 500. Per-server work is constant from 10^{11} on, two billion rows, so the step from 50 to 500 servers changes nothing per node and changes only the number of chances to draw a slow or failing one.

3 THE CORPUS IS A SEED, NOT A FILE

The corpus is defined by a seed and regenerated on demand. Each global shard draws a fresh rank-16 Gaussian basis, projects Gaussian coefficients through it with a linear decay of singular values from 1.0 to 0.3, and adds five percent noise, in an ambient dimension of 32. A shard is a pure function of its global index, so any job can produce any shard without reading anything. The companion paper [3] discusses what this corpus is and is not as a stand-in for real embeddings. For the systems question it does three things.

It removes data movement. A build from files would first have to place the corpus, 10^{12} rows of 32 single-precision values, 128 TB, where 500 jobs could read it, and then read all of it once. Built from a seed, each server’s job computes its own 400 shards and writes only its share of the 24 TB index. What moves is the code, its dependencies and the fitted bootstrap on the shared volume, which every job reads, and the query cache and partial results, which are small.

It bounds memory. A build job streams its 400 shards from the seed in row bands, writes each shard as it goes, and never holds a whole shard in memory. After each band it flushes the written pages to the volume and tells the kernel to drop them, because the page cache of freshly written data counts against the job’s two-gigabyte limit and dirty pages cannot be reclaimed until they are written. The job’s own memory stays at a few hundred megabytes, one build pod measured 292 MiB of anonymous memory beside 1526 MiB of file pages, and the build runs in the same one-CPU, two-gigabyte class as the measurement.

It makes loss recoverable. The volumes are single-replica block volumes. Six of them became unreachable when their node went offline. Each was deleted, recreated under the same name and rebuilt from its seed, on the first attempt, in 2 hours 32 minutes to 3 hours 49 minutes. No backup of the index was kept, because the seed and the code reproduce it.

Build. The 500 servers were built between 2026-08-04 and 2026-09-24. The first 16 were built in waves and the remaining 484 by one pool with no wave barrier, the same pool the measurement later used. Each server took between two and a half and four hours.

4 STORAGE IS ONE VOLUME PER SERVER

Each server’s index lives on its own block volume of 56 GiB, mounted by exactly one job at a time. Every job of a phase reads only its own volume and the shared network file system, so a slow volume slows one server and no other. Volumes are provisioned at a paced rate, because a burst of claims times out the provisioner. In August a new volume bound in about 31 seconds, and the six replacements of September took about ten minutes. Scans read the index through memory maps in blocks, so the page cache holds recently read index pages. Those pages count against the job’s memory limit, and the kernel reclaims them under it, because they are clean copies of data on the volume.

Table 2: The phases of the measurement. Every job writes its output under a temporary name and renames it, and exits at once if the output exists, so any job can be re-run at any time. The metadata phases read only partials and the index’s cell metadata.

phase	jobs	reads	writes
query cache	1	seed	100 queries
exact scan	500	own volume, queries	top-10 per query
routed pass	500	own volume, queries	top-10 per width
score	1	1500 partials	recall at ten
second sweep	500	own volume, queries	top-10 at 16, 64, 256
cell census	500	own cell metadata	cell counts
census merge	1	counts, partials	cell ranks

5 THE MEASUREMENT IS A GRAPH OF IDEMPOTENT PHASES

Table 2 lists the phases. Each is a set of independent jobs, one per server or one in all, and the driver runs the phases in order, so a phase that reads another’s output starts after it. Within a phase there is no order and no barrier, so a slow server delays only the phase’s last result. Every job reads its own volume and the shared volume and writes one small file to the shared volume. The driver checkpoints each phase to a state file, and after a restart it reloads that file and adopts the jobs still running on the cluster.

Idempotency is the property the rest of the design relies on. A job whose output exists prints that fact and exits, and a job interrupted mid-write leaves only a temporary file that the next attempt overwrites. A duplicate job, a retried job and a job the driver adopted after a restart are therefore all harmless, and the recovery rules of Section 8 can re-issue a job whenever its state is in doubt without risking a wrong or double-counted partial.

6 SUBMISSIONS YIELD TO THE CLUSTER

Jobs are not created by the driver directly. The driver publishes each job to a controller [1] that runs beside it and decides when to create it. Before every creation the controller looks at its own queue and at the cluster. It creates nothing while it already has five of its own jobs pending or twenty running, and it waits while more than a hundred pods are pending across the cluster or the target nodes are busy. A deferred job is retried after thirty seconds, then after doubling waits up to fifteen minutes, for up to fifteen attempts. The scheme is the carrier-sense and back-off rule of shared radio channels applied to a shared batch system, and it follows an earlier design for polite submission to shared clusters [2]. The effect is that the run takes capacity when the cluster has it and gives way when other groups need it.

The pool above the controller has its own rules, listed in Table 3. The one that matters most is patience with the controller. A job the controller has deferred is not yet visible to the cluster, and a driver that reads its absence as loss resubmits it and queues a duplicate behind it. The pool therefore waits ninety minutes for a submitted job to appear before it treats the job as lost.

Table 3: The submission controller’s settings during the run and the pool’s rules. The controller decides when a submitted job is created. The pool decides when a job is re-issued.

setting	value
<i>controller</i>	
own running jobs at most	20
own pending jobs at most	5
cluster pending pods that trigger waiting	100
node CPU load that triggers waiting	0.85
wait before retrying a deferred job	30 s, doubling to 15 min
attempts before giving up	15
<i>pool</i>	
jobs in flight at most	20
attempts per server	40
back-off after a failed job	2 min, doubling to 30 min
wait for a submitted job to appear	90 polls, one a minute
pending pod recycled after (first run)	45 min
terminating pod removed after (second sweep)	30 min
running job logged, not removed (second sweep)	16 h
sweeps over parked servers	12, 30 min apart

Table 4: What one job of each kind used, in the one-CPU, two-gigabyte class. Memory is what the job’s control group is charged, which includes page cache. Own memory is the process’s resident set, or its anonymous memory for the build. The build row is one of the six servers rebuilt in September. Cells the record does not hold are left empty.

job	wall	mean cores	mean mem.	peak mem.	own mem.
build	2.5–3.8 h			2048 MiB	292 MiB
exact scan	1912 s	0.88	1242 MiB	1469 MiB	346 MB
routed pass	36 min	0.48	1258 MiB	1568 MiB	165 MB

7 EVERY JOB FITS ONE CPU AND TWO GIGABYTES

Every job in the measurement requests one CPU and two gigabytes of memory. That is the smallest class the cluster offers, it is the class in which a fleet of hundreds of small jobs takes the least from other users, and the sharded index needs nothing more once three choices are made in how a scan over memory-mapped shards is written. The first is that the queries come from the cache. Deriving the 100 queries from the seed means regenerating four whole five-million-row blocks, about 640 MB each plus temporaries, so the query-cache phase writes the set once and every scan reads the file. The second is that at most two shards are open. A memory-mapped shard still materializes its row-id and tombstone arrays, 45 MB per five-million-row shard, and the sharded index holds 128 shards open by default, 5.8 GB across a full scan. The scan holds two shards open and the routed pass eight, with two worker threads. The third is that scoring proceeds in blocks of 65536 rows, which bounds the per-block temporaries to tens of megabytes.

Table 4 gives what one job of each kind used. The exact scan is compute-bound and holds nearly one core. The routed pass reads




Figure 2: The run on the calendar. The index was built over seven weeks by a pool of one job per server. The measurement, the census and the second sweep took four days at the end of September, all through the same controller.




Figure 3: Jobs the pool had in flight, from the driver logs, for the first run (exact scan, routed passes at 32 and 128 probes, score) and the second sweep (routed passes at 16, 64 and 256 probes, score). The pool holds the controller’s cap of twenty until each phase drains.

less and computes less per byte, so it averages about half a core. In both, most of the charged memory is page cache from the memory-mapped index, clean pages that the kernel reclaims under the limit, and the processes themselves hold under 350 MB. A routed pod was observed at the two-gigabyte limit with 224 MB anonymous and 1.9 GB of file pages, and it completed normally. The build reaches the limit for the same reason, with written rather than read pages, which is why it flushes and drops them band by band. The score job, which merges 1500 small partials, and the metadata jobs, which read partials and cell metadata, run in the same class in minutes and half a minute respectively.

8 A POOL OF IDEMPOTENT JOBS CARRIES THE RUN

The measurement was driven by a pool of at most twenty jobs in flight with no wave barrier, checkpointed per phase, restartable, and able to adopt jobs already on the cluster after a restart. Figure 2 places it on the calendar, Figure 3 shows how full the pool stayed, and Table 5 gives what it did. The median number of jobs in flight was nineteen in both runs. The first run finished 835 jobs that it ran itself in 39 hours, and found 14 more whose output already existed, and the second sweep finished 501 in 39 hours. The rules come first because the counts are only meaningful against them.

Rules. A job whose Kubernetes object reports success is done. A job whose object reports failure is recycled, meaning its object is deleted and, after an exponential back-off from two minutes to thirty, resubmitted, up to forty times before the server is parked and swept again later. A job submitted through the controller but

Table 5: The pool over the first run, 2026-09-25 18:01 UTC to 2026-09-27 08:59 UTC, and the second sweep, 2026-09-27 21:37 UTC to 2026-09-29 12:56 UTC. First-run scan counts exclude 152 servers scanned in an earlier session at a larger job shape.

	first run	second sweep
submissions	928	574
completions	849	501
recycles after a failed job	75	73
storage read fault on one node	24	25
node lost during the job	13	17
node unreachable, pod stuck	13	11
container start failure	2	1
other	23	19
recycles after 45 minutes pending	26	0
submissions deferred by the controller	1	0
servers parked or given up	0	0

not yet visible as an object is waited on for ninety polls, one a minute, before it is treated as lost. In the first run a job whose pod had not reached the running state 45 minutes after submission was recycled the same way. For the second sweep that rule was replaced by waiting, because a running job is never deleted on a namespace other groups share, and a job running longer than sixteen hours is logged rather than removed.

The deferral rule. The submission controller defers creating a job while the namespace already has several pending or the cluster is busy, and retries for up to about an hour. A driver that treats a deferred submission as a lost one resubmits, and every resubmission queues a duplicate behind the original it was meant to replace. The wait before a submission is treated as lost separates the two cases. The run needed the rule once, for the score job.

Failure classes. Two dominated, and both are ordinary on a shared cluster. One node returned a storage read fault to every job placed on it within 3 to 13 minutes of start, after its block volume had attached. Retries landed elsewhere and finished, at a cost of about 130 slot-minutes out of 9600 over the first eight hours. The second class, pods pending past 45 minutes, was benign in every instance checked, because the original pod had run and the retry found its output within minutes. That is what idempotent jobs buy. In the second sweep 63 servers needed a second attempt, 9 a third and 1 a fourth, and a storage fault at one moment ended all twenty jobs in flight at once, after which fresh submissions ran and the batch was re-issued after its back-off.

Pods that stop answering. A pod can stay in a terminating state on a node that no longer responds, or run on a node whose storage path is slow, and the job controller does not replace it on its own. Four such pods in the first run and two in the second sweep were removed by hand, and their replacements finished normally. The second sweep added a rule for pods terminating longer than thirty minutes. Across four days the manual actions were six pod removals and one resubmission of the score job.

Table 6: Per-server wall time in the 10^{12} run, one CPU per job, 500 servers per phase, and the CPU-hours each phase consumed. The three routed widths of the second sweep ran inside one job per server and share its start-up. Each time includes pulling the image, installing dependencies, cloning the repository and opening 400 shards from network block storage, so it is the cost of the measurement job and not of inverted-file search. Derived from the record by `paper/pvldb1t/figdata/fig_data.py`.

phase	median	mean	p10	p90	max	CPU-h
exact scan	2664 s	3008 s	1850 s	4009 s	16308 s	418
routed, 32	1064 s	1131 s	775 s	1466 s	14286 s	157
routed, 128	1335 s	1420 s	958 s	1861 s	14524 s	197
routed, 16	916 s	1071 s	660 s	1440 s	12236 s	149
routed, 64	1219 s	1366 s	912 s	1743 s	14645 s	190
routed, 256	1594 s	1719 s	1176 s	2199 s	13786 s	239




Figure 4: Distribution of per-server wall time by phase, 500 servers each, one CPU per job. The bodies are tight and the tails are a few servers on slow nodes, which the pool absorbs by waiting rather than by re-placing. Data derived from the record logs.

Provenance. The job containers cloned the repository’s default branch at run time, and twelve commits appear in pod logs across the first run. The phase scripts came from a configuration map fixed for each run, and the merge and score code did not change. Pinning the clone to a commit is the one change to the descriptor for the next run.

9 THE MEASUREMENT COST 1349 CPU-HOURS

Table 6 and Figure 4 give what the measurement cost. The exact scan of two billion rows at one CPU takes 44 minutes at the median and the routed pass 15 to 27 minutes depending on width. The first run, exact scan and routed passes at 32 and 128, consumed 772 CPU-hours, and the second sweep at 16, 64 and 256 consumed 577, so the whole measurement is 1349 CPU-hours of the cluster’s smallest job class. The routed times are not the cost of inverted-file search, and the ratio between the routed and exact rows is not a latency figure for the method. These jobs are read-bound on network block storage, spend a fixed minute or two on image pull,

dependency install and clone, and open 400 shards before scoring. The 10^{11} run, on six CPUs per job, reported 14.7 and 8.5 as the same ratio, and neither number is the method’s. The tails are node variance. Between 4 and 30 servers per phase took more than twice the median, the 16308 s maximum of the scan is one server on one node that also produced the slowest routed pass, and in every phase the ninetyth percentile is within 1.6 times the median.

10 SEVEN DESIGN CHOICES, EACH AGAINST ITS OBVIOUS ALTERNATIVE

Table 1 collects the choices. Each one replaces a choice that is reasonable at a billion rows with one that stays reasonable at a trillion, and each is measured in the run rather than argued. None depends on vector search. A seeded generator, bounded jobs that are safe to repeat, and a pool that waits for the platform instead of racing it apply to any computation that is too large for one machine and too small to justify reserved capacity.

We would change three things. The job containers cloned the repository’s default branch at run time, so the code that ran moved as the branch moved, and the next run pins the clone to one commit. A rule for pods that stay in a terminating state now exists in the driver, and the next run starts with it. The query set came from four seeded shards, which concentrates the reference neighbours on few servers, and the companion paper measures a query set drawn from seeds no corpus shard uses.

11 RELATED WORK

Distributed vector databases such as Milvus [11] and its cloud-native successor [6] shard an index across nodes and merge per-node results, which is the arrangement our fleet uses without a serving layer. Billion-scale practice with Faiss [5, 7], DiskANN [10] and SPANN [4] places an index on one machine or one node’s disk, and the Big ANN competitions [8, 9] fixed datasets and reporting at 10^9 rows. Recent work places a billion-scale index on distributed storage [13]. The National Research Platform [12] describes the federated cluster this run used. We found no published account of building and exactly scanning a 10^{12} -row vector index on a shared cluster in its smallest job class.

12 REPRODUCIBILITY

The repository at the availability URL carries the corpus seeds and generator, the build, scan, routed, score and metadata scripts, the pool driver with its rules, and, under the directory `benchmarks/fleet/record/1t/`, the pool state files, every driver log, the pilot logs and the score logs whose `RESULT_JSON` lines hold every per-server wall time. The cost table and the wall-time figure are produced from those logs by `paper/pvldb1t/figdata/fig_data.py`. Rebuilding the index from seed costs about three and a half hours per server at the pool’s width.

13 CONCLUSION

A trillion-row vector index was built from a seed and measured exactly on a shared research cluster in the smallest job class it offers. The build moved no corpus bytes, a lost volume was a four-hour rebuild, and a pool of idempotent jobs with explicit rules for failure,

waiting and deferral carried the four-day measurement, 1349 CPU-hours, through the failures a shared cluster produces, with six pods replaced by hand.

ACKNOWLEDGMENTS

This work used the National Research Platform’s Nautilus cluster. The National Research Platform is supported in part by National Science Foundation awards CNS-1730158, ACI-1540112, ACI-1541349, OAC-1826967, OAC-2112167, CNS-2100237 and CNS-2120019. We thank its operators and community.

REFERENCES

- [1] Andrew H. Bond. 2026. nats-bursting: submitting batch jobs to a shared Kubernetes cluster over a NATS bus, with politeness back-off. <https://github.com/ahb-sjsu/nats-bursting>.
- [2] Andrew H. Bond. 2026. polite-submit: client-side contention management for Slurm clusters with carrier-sense back-off. <https://github.com/ahb-sjsu/polite-submit>.
- [3] Andrew H. Bond. 2026. Recall Does Not Move at a Trillion Rows [Experiment, Analysis & Benchmark]. Under submission to PVLDB.
- [4] Qi Chen, Bing Zhao, Haidong Wang, Mingqin Li, Chuanjie Liu, Zengzhong Li, Mao Yang, and Jingdong Wang. 2021. SPANN: Highly-Efficient Billion-Scale Approximate Nearest Neighbor Search. In *Advances in Neural Information Processing Systems (NeurIPS)*.
- [5] Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvasy, Pierre-Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. 2024. The Faiss Library. *arXiv preprint arXiv:2401.08281* (2024).
- [6] Rentong Guo, Xiaofan Luan, Long Xiang, Xiao Yan, Xiaomeng Yi, Jigao Luo, Qianya Cheng, Weizhi Xu, Jiarui Luo, Frank Liu, Zhenshan Cao, Yanliang Qiao, Ting Wang, Bo Tang, and Charles Xie. 2022. Manu: A Cloud Native Vector Database Management System. *Proceedings of the VLDB Endowment* 15, 12 (2022), 3548–3561. <https://doi.org/10.14778/3554821.3554843>
- [7] Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2021. Billion-Scale Similarity Search with GPUs. *IEEE Transactions on Big Data* 7, 3 (2021), 535–547.
- [8] Harsha Vardhan Simhadri, Martin Aumüller, Amir Ingber, Matthijs Douze, George Williams, Magdalen Dobson Manohar, Dmitry Baranchuk, Edo Liberty, Frank Liu, Ben Landrum, Mazin Karjkar, Laxman Dhulipala, Meng Chen, Yue Chen, Rui Ma, Kai Zhang, Yuzheng Cai, Jiayang Shi, Yizhuo Chen, Weiguo Zheng, Zihao Wan, Jie Yin, and Ben Huang. 2024. Results of the Big ANN: NeurIPS’23 competition. *arXiv preprint arXiv:2409.17424* (2024).
- [9] Harsha Vardhan Simhadri, George Williams, Martin Aumüller, Matthijs Douze, Artem Babenko, Dmitry Baranchuk, Qi Chen, Lucas Hosseini, Ravishankar Krishnaswamy, Gopal Srinivasa, Suhas Jayaram Subramanya, and Jingdong Wang. 2022. Results of the NeurIPS’21 Challenge on Billion-Scale Approximate Nearest Neighbor Search. In *Proceedings of the NeurIPS 2021 Competitions and Demonstrations Track (Proceedings of Machine Learning Research)*, Vol. 176. 177–189.
- [10] Suhas Jayaram Subramanya, Devvrit, Rohan Kadekodi, Ravishankar Krishnaswamy, and Harsha Vardhan Simhadri. 2019. DiskANN: Fast Accurate Billion-point Nearest Neighbor Search on a Single Node. In *Advances in Neural Information Processing Systems*, Vol. 32.
- [11] Jianguo Wang, Xiaomeng Yi, Rentong Guo, Hai Jin, Peng Xu, Shengjun Li, Xiangyu Wang, Xiangzhou Guo, Chengming Li, Xiaohai Xu, Kun Yu, Yuxing Yuan, Yinghao Zou, Jiquan Long, Yudong Cai, Zhenxiang Li, Zhifeng Zhang, Yihua Mo, Jun Gu, Ruiyi Jiang, Yi Wei, and Charles Xie. 2021. Milvus: A Purpose-Built Vector Data Management System. In *Proceedings of the 2021 International Conference on Management of Data (SIGMOD)*. 2614–2627. <https://doi.org/10.1145/3448016.3457550>
- [12] Derek Weitzel, Ashton Graves, Sam Albin, Huijun Zhu, Frank Würthwein, Mahidhar Tatineni, Dmitry Mishin, John Graham, Elham E. Khoda, Mohammad Firas Sada, Larry Smarr, and Thomas DeFanti. 2025. The National Research Platform: Stretched, Multi-Tenant, Scientific Kubernetes Cluster. In *Practice and Experience in Advanced Research Computing (PEARC ’25)*. arXiv:2505.22864.
- [13] Kun Yu, Jiabao Jin, Xiaoyao Zhong, Peng Cheng, Lei Chen, Zhitao Shen, Jingkuan Song, Hengtao Shen, and Xuemin Lin. 2025. Approximate Nearest Neighbor Search of Large Scale Vectors on Distributed Storage. *arXiv preprint arXiv:2510.17326* (2025).